

Package ‘STITCH’

June 21, 2018

Type Package

Title STITCH - Sequencing To Imputation Through Constructing Haplotypes

Version 1.5.1

Date 2017-06-21

Author Robert William Davies

Maintainer Robert William Davies <robertwilliamdavies@gmail.com>

Description STITCH performs imputation of individuals sequenced to low coverage in a read aware fashion without a reference panel.

Installation To install, first install dependencies, then run the `install.packages` command, pointing to the downloaded tarball (STITCH.tar.gz)

Getting started A minimum run requires the following options to be set: the chromosome being run (`chr`); a path to a file with a set of bi-allelic SNP sites (`posfile`); a choice of `K`, the number of internally modelled haplotypes (`K`); a path to an output directory (`outputdir`); a path to a temporary directory, ideally on fast disks or a RAM disk (`tempdir`); a list of bam files (`bamlist`); and the number of generations since founding (`nGen`), which can be approximated from a choice of `K` for wild populations from $4 * Ne / K$. Additional useful options relate to what region to impute (`regionStart`, `regionEnd`, `buffer`), whether to use validation data to benchmark imputation (`genfile`), the number of cores to use (`nCores`), whether imputation is run on a server or cluster (`environment`), the number of EM iterations (`niterations`), whether to run in diploid or pseudoHaploid mode (`method`), and if run in pseudoHaploid mode, what iteration to switch from pseudoHaploid to diploid (`switchModelIteration`).

Depends parallel, rrbgen (>= 0.0.1)

Imports Rcpp, data.table

LinkingTo Rcpp, RcppArmadillo

RoxygenNote 6.0.1

License GPL | file LICENSE

SystemRequirements C++11

NeedsCompilation yes

Suggests testthat

R topics documented:

make_STITCH_cli	2
STITCH	2
Index	8

make_STITCH_cli	<i>Make STITCH command line interface</i>
-----------------	---

Usage

```
make_STITCH_cli(function_file, cli_output_file,
  integer_vectors = c("shuffleHaplotypeIterations", "splitReadIterations",
    "refillIterations"), character_vectors = c("reference_populations"),
  other_logical_params = NULL, other_integer_params = NULL,
  other_double_params = NULL, other_character_params = NULL,
  function_name = "STITCH")
```

Arguments

function_file to main STITCH function file
stitch_cli_file where output goes

STITCH	<i>Sequencing To Imputation Through Constructing Haplotypes</i>
--------	---

Usage

```
STITCH(chr, nGen, posfile, K, outputdir, tempdir = NA, bamlist = "",
  cramlist = "", reference = "", genfile = "", method = "diploid",
  output_format = "bgvcf", B_bit_prob = 16,
  outputInputInVCFFormat = FALSE, downsampleToCov = 50,
  downsampleFraction = 1, readAware = TRUE, chrStart = NA, chrEnd = NA,
  regionStart = NA, regionEnd = NA, buffer = NA,
  maxDifferenceBetweenReads = 1000, maxEmissionMatrixDifference = 1e+10,
  alphaMatThreshold = 1e-04, emissionThreshold = 1e-04,
  iSizeUpperLimit = as.integer(600), bqFilter = as.integer(17),
  niterations = 40, shuffleHaplotypeIterations = c(4, 8, 12, 16),
  splitReadIterations = 25, nCores = 1, expRate = 0.5, maxRate = 100,
  minRate = 0.1, Jmax = 1000, regenerateInput = TRUE,
  originalRegionName = NA, keepInterimFiles = FALSE, keepTempDir = FALSE,
  environment = "server", outputHaplotypeProbabilities = FALSE,
  switchModelIteration = NA, generateInputOnly = FALSE,
  restartIterations = NA, refillIterations = c(6, 10, 14, 18),
  downsampleSamples = 1, downsampleSamplesKeepList = NA,
  subsetSNPsfile = NA, useSoftClippedBases = FALSE,
  outputBlockSize = 1000, inputBundleBlockSize = NA,
  reference_haplotype_file = "", reference_legend_file = "",
```

```
reference_sample_file = "", reference_populations = NA,
reference_phred = 20, reference_iterations = 10, output_filename = NULL,
initial_min_hapProb = 0.4, initial_max_hapProb = 0.6,
regenerateInputWithDefaultValues = FALSE,
plotHapSumDuringIterations = FALSE,
plot_shuffle_haploptype_attempts = FALSE, plotAfterImputation = TRUE,
save_sampleReadsInfo = FALSE, gridWindowSize = NA,
shuffle_bin_nSNPs = NULL, shuffle_bin_radius = 5000)
```

Arguments

chr	What chromosome to run. Should match BAM headers
nGen	Number of generations since founding or mixing. Note that the algorithm is relatively robust to this. Use $nGen = 4 * Ne / K$ if unsure
posfile	Where to find file with positions to run. File is tab separated with no header, one row per SNP, with col 1 = chromosome, col 2 = physical position (sorted from smallest to largest), col 3 = reference base, col 4 = alternate base. Bases are capitalized. Example first row: 1<tab>1000<tab>A<tab>G<tab>
K	How many founder / mosaic haplotypes to use
outputdir	What output directory to use
tempdir	What directory to use as temporary directory. If set to NA, use default R tempdir. If possible, use ramdisk, like /dev/shm/
bamlist	Path to file with bam file locations. File is one row per entry, path to bam files. Bam index files should exist in same directory as for each bam, suffixed either .bam.bai or .bai
cramlist	Path to file with cram file locations. File is one row per entry, path to cram files. cram files are converted to bam files on the fly for parsing into STITCH
reference	Path to reference fasta used for making cram files. Only required if cramlist is defined
genfile	Path to gen file with high coverage results. Empty for no genfile. File has a header row with a name for each sample, matching what is found in the bam file. Each subject is then a tab separated column, with 0 = hom ref, 1 = het, 2 = hom alt and NA indicating missing genotype, with rows corresponding to rows of the posfile. Note therefore this file has one more row than posfile which has no header
method	How to run imputation - either diploid, pseudoHaploid, or diploid-inbred. Please see main README for more information. All methods assume diploid samples. diploid is the most accurate but slowest, while pseudoHaploid may be advantageous for large sample sizes and K. diploid-inbred assumes all samples are inbred and invokes an internal haploid mathematical model but outputs diploid genotypes and probabilities
output_format	one of bgvcf (i.e. bgzipped VCF) or bgen (Layout = 2, CompressedSNPBlocks = 1)
B_bit_prob	when using bgen, how many bits to use to store each double. Options are 8, 16, 24 or 32
outputInputInVCFFormat	Whether to output the input in vcf format
downsampleToCov	What coverage to downsample individual sites to. This ensures no floating point errors at sites with really high coverage

downsampleFraction	Downsample BAMs by choosing a fraction of reads to retain. Must be value $0 < \text{downsampleFraction} < 1$
readAware	Whether to run the algorithm in read aware mode. If false, then reads are split into new reads, one per SNP per read
chrStart	When loading from BAM, some start position, before SNPs occur. Default NA will infer this from either regionStart, regionEnd and buffer, or posfile
chrEnd	When loading from BAM, some end position, after SNPs occur. Default NA will infer this from either regionStart, regionEnd and buffer, or posfile
regionStart	When running imputation, where to start from. The 1-based position x is kept if $\text{regionStart} \leq x \leq \text{regionEnd}$
regionEnd	When running imputation, where to stop.
buffer	Buffer of region to perform imputation over. So imputation is run from regionStart+buffer to regionEnd+buffer, and reported for regionStart to regionEnd, including the bases of regionStart and regionEnd
maxDifferenceBetweenReads	How much of a difference to allow the reads to make in the forward backward probability calculation. For example, if $P(\text{read} \mid \text{state } 1) = 1$ and $P(\text{read} \mid \text{state } 2) = 1e-6$, re-scale so that their ratio is this value. This helps prevent any individual read as having too much of an influence on state changes, helping prevent against influence by false positive SNPs
maxEmissionMatrixDifference	Similar to maxDifferenceBetweenReads, specifies ratio of how much larger the most probable state can be than the least probable state, but across all reads rather than for a single read. This helps to limit overflow in C++ calculations
alphaMatThreshold	Minimum (maximum is 1 minus this) state switching into probabilities
emissionThreshold	Emission probability bounds. $\text{emissionThreshold} < P(\text{alt read} \mid \text{state } k) < (1 - \text{emissionThreshold})$
iSizeUpperLimit	Do not use reads with an insert size of more than this value
bqFilter	Minimum BQ for a SNP in a read. Also, the algorithm uses $\text{bq} \leq \text{mq}$, so if mapping quality is less than this, the read isn't used
niterations	Number of EM iterations.
shuffleHaplotypeIterations	Iterations on which to perform heuristic attempt to shuffle founder haplotypes for better fit. To disable set to NA.
splitReadIterations	Iterations to try and split reads which may span recombination breakpoints for a better fit
nCores	How many cores to use
expRate	Expected recombination rate in cM/Mb
maxRate	Maximum recomb rate cM/Mb
minRate	Minimum recomb rate cM/Mb
Jmax	Maximum number of SNPs on a read

regenerateInput	Whether to regenerate input files. If this is FALSE, please using the same region-Start, regionEnd, buffer and posfile as you used to generate the input. Setting any of those to different values can cause the previous input data to be improperly interpreted. Please also see originalRegionName and regenerateInputWithDefaultValues
originalRegionName	If regenerateInput is FALSE (i.e. using existing data), this is the name of the original region name (chr.regionStart.regionEnd). This is necessary to load past variables
keepInterimFiles	Whether to keep interim parameter estimates
keepTempDir	Whether to keep files in temporary directory
environment	Whether to use server or cluster multicore options
switchModelIteration	Whether to switch from pseudoHaploid to diploid and at what iteration (NA for no switching)
generateInputOnly	Whether to just generate input data then quit
restartIterations	In pseudoHaploid method, which iterations to look for collapsed haplotype probabilities to resolve
refillIterations	When to try and refill some of the less frequently used haplotypes
downsampleSamples	What fraction of samples to retain. Useful for checking effect of N on imputation. Not meant for general use
downsampleSamplesKeepList	When downsampling samples, specify a numeric list of samples to keep
subsetSNPsfile	If input data has already been made for a region, then subset down to a new set of SNPs, as given by this file. Not meant for general use
useSoftClippedBases	Whether to use (TRUE) or not use (FALSE) bases in soft clipped portions of reads
outputBlockSize	How many samples to write out to disk at the same time when making temporary VCFs that are later pasted together at the end to make the final VCF. Smaller means lower RAM footprint, larger means faster write.
inputBundleBlockSize	If NA, disable bundling of input files. If not NA, bundle together input files in sets of <= inputBundleBlockSize together
reference_haplotype_file	Path to reference haplotype file in IMPUTE format (file with no header and no rownames, one row per SNP, one column per reference haplotype, space separated, values must be 0 or 1)
reference_legend_file	Path to reference haplotype legend file in IMPUTE format (file with one row per SNP, and a header including position for the physical position in 1 based coordinates, a0 for the reference allele, and a1 for the alternate allele)

reference_sample_file	Path to reference sample file (file with header, one must be POP, corresponding to populations that can be specified using reference_populations)
reference_populations	Vector with character populations to include from reference_sample_file e.g. CHB, CHS
reference_phred	Phred scaled likelihood or an error of reference haplotype. Higher means more confidence in reference haplotype genotypes, lower means less confidence
reference_iterations	When using reference haplotypes, how many iterations to use to train the starting data
output_filename	Override the default bgzip-VCF / bgen output name with this given file name. Please note that this does not change the names of inputs or outputs (e.g. RData, plots), so if outputdir is unchanged and if multiple STITCH runs are processing on the same region then they may over-write each others inputs and outputs
initial_min_hapProb	Initial lower bound for probability read comes from haplotype. Double bounded between 0 and 1
initial_max_hapProb	Initial upper bound for probability read comes from haplotype. Double bounded between 0 and 1
regenerateInputWithDefaultValues	If regenerateInput is FALSE and the original input data was made using region-Start, regionEnd and buffer as default values, set this equal to TRUE
plotHapSumDuringIterations	Boolean TRUE/FALSE about whether to make a plot that shows the relative number of individuals using each ancestral haplotype in each iteration
plot_shuffle_haplotype_attempts	Boolean TRUE/FALSE about whether to make a plot that tries to show the selection of ancestral haplotypes to check for shuffling / flipping
plotAfterImputation	Boolean TRUE/FALSE about whether to make plots after imputation has run (can be set to FALSE if this throws errors on systems without x11)
save_sampleReadsInfo	Experimental. Boolean TRUE/FALSE about whether to save additional information about the reads that were extracted
gridWindowSize	Whether to work on a grid where reads are binned into windows of this size (1 based, i.e. first bin is bases 1-gridWindowSize). This is particularly appropriate for very low coverage data (e.g. less than 0.2X) and can substantially speed up analyses
shuffle_bin_nSNPs	Parameter that controls how to detect ancestral haplotypes that are shuffled during EM for possible re-setting. If set (not NULL), then break per-SNP (or per-grid) every this many SNPs / grids, and compare each to detect whether haplotypes either 1) are more likely to stay where they are or 2) switch from one haplotype to another. Note that only one of shuffle_bin_nSNPs or shuffle_bin_radius should be non-NULL

shuffle_bin_radius

Parameter that controls how to detect ancestral haplotypes that are shuffled during EM for possible re-setting. If set (not NULL), then recombination rate is calculated around pairs of SNPs in window of twice this value, and those that exceed what should be the maximum (defined by nGen and maxRate) are checked for whether they are shuffled

Value

Results in properly formatted version

Author(s)

Robert Davies

Index

`make_STITCH_cli`, [2](#)

`STITCH`, [2](#)